

Sparse Retrieval in the Age of RAG

Antonio Mallia

LiveRAG Challenge @ SIGIR'2025

July 17, 2025

About me



✉ me@antoniomallia.it
🏠 www.antoniomallia.it

X @antonio_mallia
Cat amallia
in in/antoniomallia

I am currently a **Staff Research Scientist** at Pinecone.

Prior to this, I served as an Applied Scientist on the Artificial General Intelligence (AGI) team at Amazon.

I hold a Ph.D. from New York University, where my research focused on efficient web retrieval methodologies.



Pinecone

Founded by Edo Liberty

Created the vector database category

Pinecone's mission is to make AI knowledgeable

Pinecone is the leading vector database for building accurate and performant AI applications at scale in production.

Go and check it out: <https://www.pinecone.io/>

Agenda

Traditional IR vs. RAG systems

Sparse Retrieval in LiveRAG

Advantages of Sparse

Not a replacement for Dense: hybrid

Future: Learned Sparse Retrieval for RAG

The Evolution of Retrieval

From Ranked Lists to Generative Evidence

Traditional IR: Ranked lists optimized for precision, recall, and user clicks.

RAG: Needs evidence retrieval that supports generation, not navigation.

What matters now:

- Faithfulness, coverage, reasoning chains
- Retrieval that supports answer synthesis, not just result finding



RAG Breaks the Old Rules

In RAG, the retriever isn't helping a user – it's helping a model. That changes everything.

Traditional IR



User browses ranked list



Optimize for top-k relevance (MAP, NDCG)



Queries are short, keyword-based



Focus on document ranking

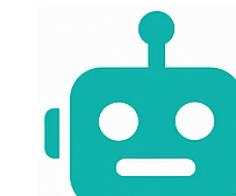


Evaluation: clicks, rank positions



Retrieval ends with the user

RAG Systems



LLM consumes retrieved content



Optimize for grounding and synthesis



Queries are natural, long-form, ambiguous



Focus on passage evidence selection




Evaluation: faithfulness, hallucination rate




Retrieval fuels generation

First Stage Retrieval

Dense Retrieval



Sparse Retrieval



T ₁	capital	→ [3]
T ₂	city	→ [3]
T ₃	conference	→ [1, 6]
T ₄	cool	→ [1, 5]
T ₅	engine	→ [5, 6]
T ₆	essir	→ [2]
T ₇	france	→ [3]
T ₈	milan	→ [2, 6]
T ₉	paris	→ [3, 4]
T ₁₀	school	→ [2]
T ₁₁	search	→ [5, 6]
T ₁₂	sigir	→ [1, 4]
T ₁₃	summer	→ [2]
T ₁₄	year	→ [4]

Terms = {capital, city, conference, cool, engine, essir, france, milan, paris, school, search, sigir, summer, year}

Sparse Retrieval Ubiquitous at LiveRAG

Nearly all* LiveRAG teams used some form of sparse in their pipeline



Most teams used BM25

CIIR (used Lion)

DoTA-RAG (prune with BM25)

GraphRAG used BM25 to seed a graph

Ragmatazz indexed with two BM25

* 17 out of 19

Sparse Retrieval @ LiveRAG

Ragmatazz

Rank	bm25s	bm25s_hyde	arc-embed	arc-embed_hyde
1	.216	.137	.15	.133
3	.341	.236	.267	.231
5	.396	.284	.337	.28
10	.472	.365	.424	.364
20	.546	.458	.518	.454
40	.616	.532	.615	.541
100	.71	.635	.713	.651
1000	.877	.844	.882	.853

Table 2: Mean Recall @ Rank for the 4 Retrievers

RAGentA

Table 2: Evaluation of Retrieval Performance

Retrieval System	MRR@20	Recall@20
BM25	0.4205	0.5020
E5	0.3476	0.4920
Hybrid	0.4290	0.5650

PreQRAG

Single-doc Questions				
	Top1	Top2	Top3	Top 10
Sparse (Rewritten)	37.5%	47.7%	52.1%	67.4%
Dense (Rewritten)	31.2%	36.2%	42.0%	52.1%
Multi-doc Questions				
	Top1	Top2	Top3	Top 10
Sparse (Rewritten)	32%	36%	40%	55%
Dense (Rewritten)	26%	30%	36%	56%

Table 2: Retrieval Performance Metrics for Single-doc and Multi-doc Questions

TopClustRAG

TABLE I

RETRIEVAL PERFORMANCE OF SPARSE, DENSE, AND HYBRID SYSTEMS ON THE SYNTHETIC VALIDATION SET.

System	MRR	R@1	R@5	R@10	R@50	R@100	R@200	R@1000
Sparse	0.3361	0.2037	0.4074	0.4815	0.5741	0.6481	0.7593	0.8704
Dense	0.0526	0.0000	0.0926	0.1111	0.2963	0.3333	0.3519	0.5556
Hybrid	0.1322	0.0370	0.1111	0.3519	0.6852	0.7778	0.8519	0.8889

Team Marikarp

Retriever	R@10	R@100	R@200	R@400	R@1k	R@2k	R@4k
sparse&dense	0.61	0.81	0.86	0.89	0.93	0.94	0.95
dense	0.53	0.72	0.77	0.84	0.88	0.90	0.93
sparse	0.59	0.76	0.82	0.86	0.89	0.93	0.94

Table 2: Recall of gold documents under single-document dataset using different retrieval methods.


DoTA-RAG

Method	Correctness [-1:2]		Faithfulness [-1:1]	
	All words	300-word cap	All words	300-word cap
Baseline	0.752	0.761	-0.496	-0.493
+ Arctic-M	1.616	1.626	-0.216	-0.225
+ Routing	1.562	1.577	-0.108	-0.090
+ Pruning	1.562	1.566	0.428	0.404
+ Rerank	1.652	1.686	0.672	0.662
+ Rewrite	1.478	1.484	0.640	0.620

Emorag

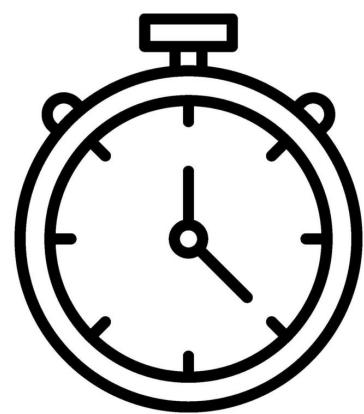
Name	MAP	Recip. Rank	nDCG@10	Recall@1	Recall@10
Sparse (OpenSearch BM25)	.523	.347	.497	.285	.485
Dense (Pinecone E5)	.352	.260	.367	.190	.435
Hybrid	.523	.347	.497	.285	.485

Efficiency




Memory footprint

Chen et al. (2022, EMNLP Findings) note that a BM25 inverted index for the MS MARCO passage corpus occupies only about 0.7 GB, whereas a dense retriever's vector index can require tens of gigabytes (approximately 26 GB in their example)



Latency and Query Throughput

Lin (2024, arXiv:2409.06464v1) shows that on the 15-million-document corpus a BM25 inverted index handles ~210 queries per second compared with ~56 QPS for a strong dense HNSW retriever, demonstrating that sparse retrieval can be faster than dense retrieval at large scale.



Compute Requirements and Deployment Cost


Lassance & Clinchant (2022) emphasize, “multi-CPU core + GPU (vs. mono-CPU core)” setups are often “the norm” for dense retrievers

Not a replacement for dense!

Sparse retrieval captures lexical overlap: precise, interpretable, efficient.

Dense retrieval captures semantic similarity: fuzzy matching, generalization.

Combined they consistently outperform either one alone.



How to combine dense and sparse?

Score Fusion

Combine the scores returned by the sparse and dense retrievers.

[RAGentA](#) [NoobRAG](#)

Alternating merge

Alternately selecting from sparse or dense

[Team Marikarp](#)

Reciprocal Rank Fusion (RRF)

Instead of scores, use rank positions

[RMIT-ADM+S](#) [TopClustRAG](#) [UiS-IAI](#) [Ragmatazz](#) [PRMAS-DRCA](#)


Retrieve + Rerank

Retrieve w/ sparse and dense → then use a reranker to find common ranking

Retrieve + Prune

Retrieve w/ dense → Sparse-based Pruning

[DoTA-RAG](#)



Hybrid and Sparse Retrieval Drive Better Coverage

“Disabling BM25 reduces Recall@5 by thirteen points because dense retrieval alone still struggles with misspellings and rare entities” - TinyUPR

Table 2: Evaluation of Retrieval Performance

Retrieval System	MRR@20	Recall@20
BM25	0.4205	0.5020
E5	0.3476	0.4920
Hybrid	0.4290	0.5650

“Our hybrid retrieval system achieves an MRR@20 of 0.4290, outperforming BM25 (0.4205) by +2.0% and E5 (0.3476) by +23.4%. ” - RAGentA



Complementarity Objectives for Hybrid Retrieval

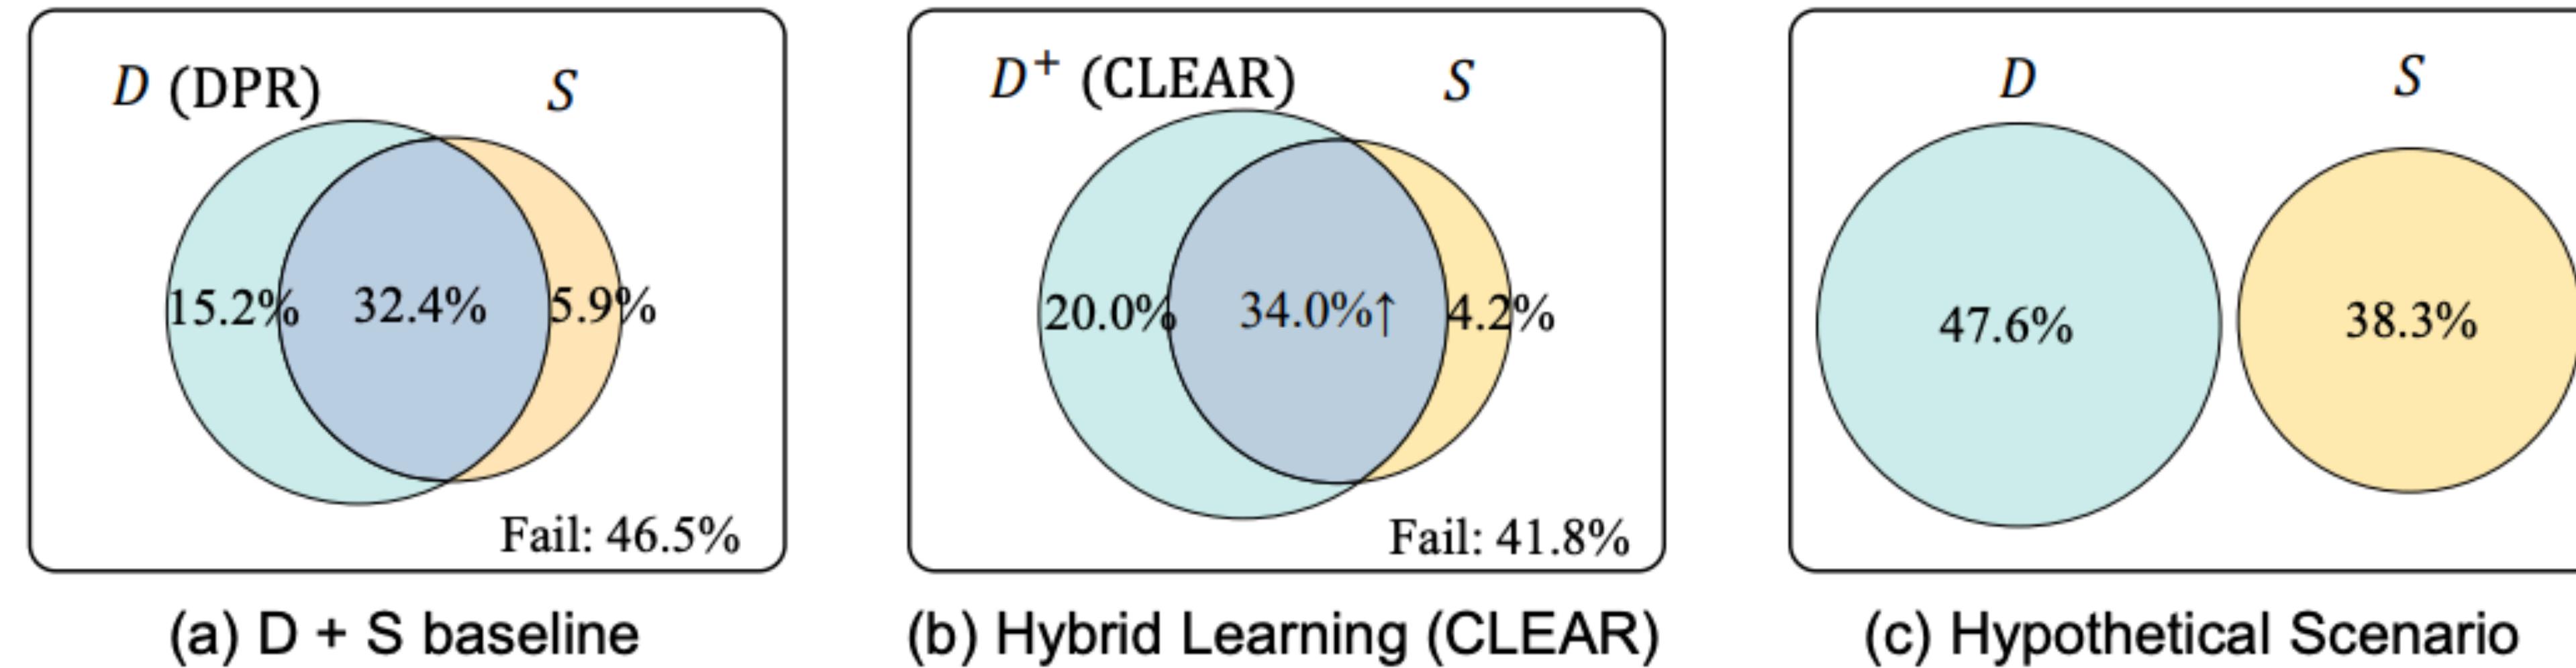


Figure 1: Recall@10 on Natural Questions. In the venn diagram, (a) shows BM25+DPR baseline and (b) shows CLEAR using residual margin. (c) is a hypothetical scenario, identical to (a) but without the intersection


What most teams use vs. what's available

Despite sparse retrieval being everywhere, few teams moved beyond BM25.

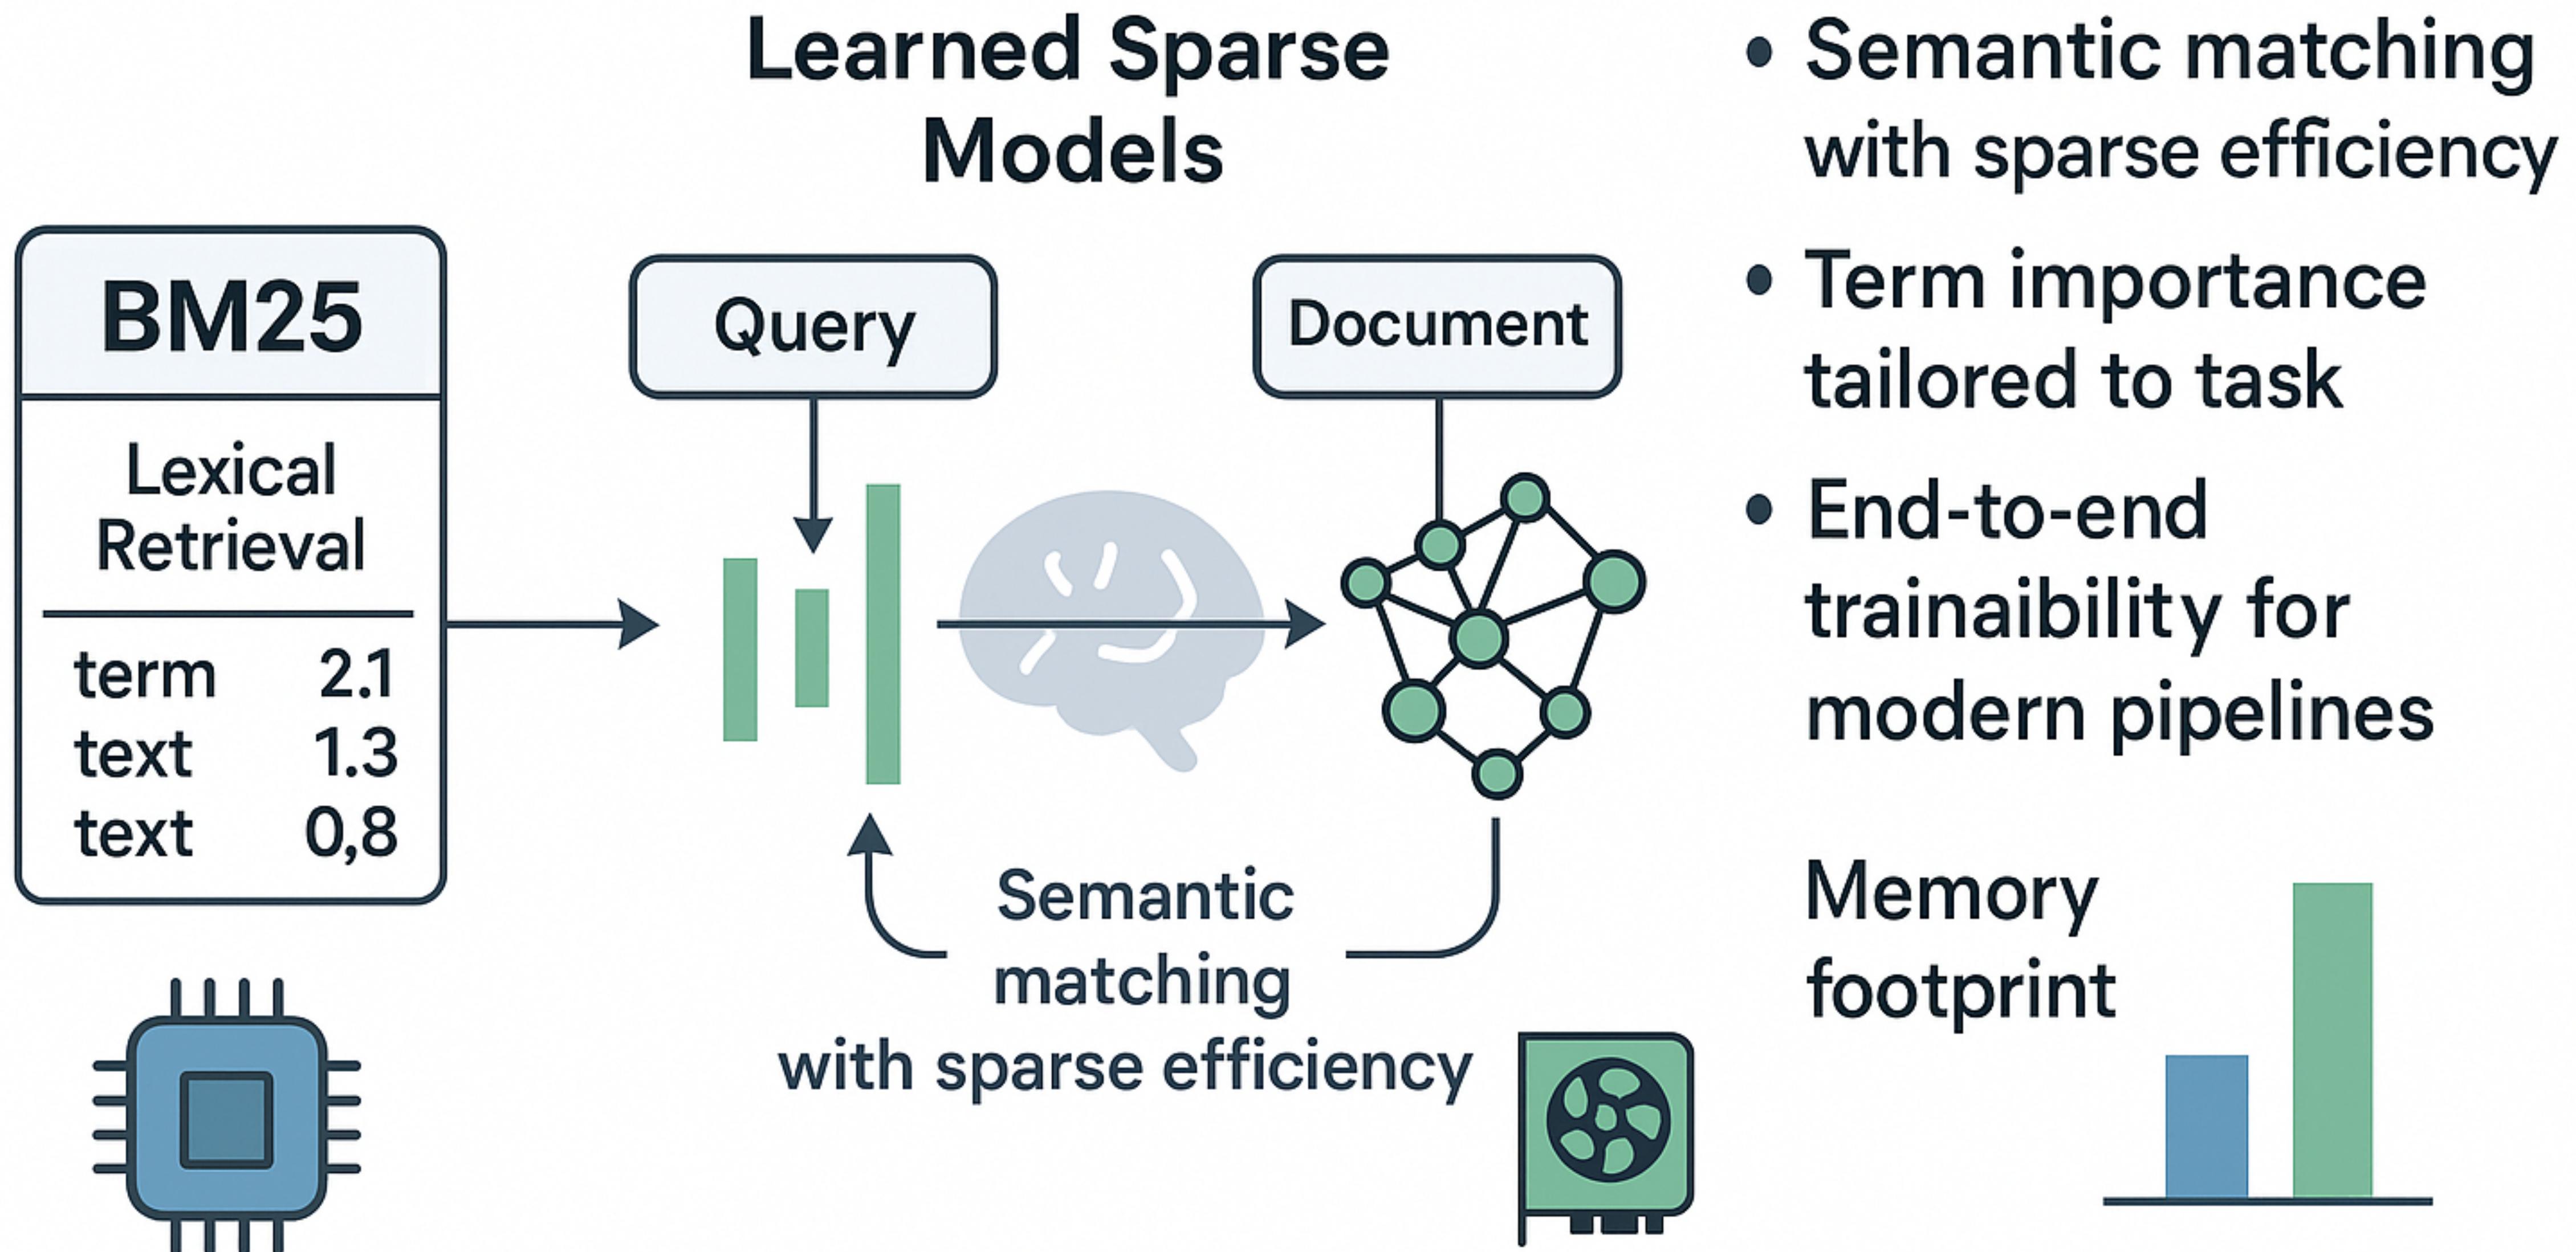
It's a decades-old method, fast and solid, but doesn't reflect recent progress in sparse neural IR.



History of learned sparse models



Enter Learned Sparse Models





An example

Predict scores for each unique term of the document

Directly store quantized scores in inverted index

Score is the sum of intersection terms

Goal: maximize the score difference between positive and negative document




Training a sparse model

Triples sampled from the MS-MARCO training dataset

A query, a relevant passage, and a presumed non-relevant passage per sample

Two scores for the corresponding two documents are computed

Pairwise softmax cross-entropy loss over the document scores



Pinecone Sparse Model

Built on top of the DeepImpact architecture

The model directly estimates the lexical importance of tokens by leveraging their context

Outperforms BM25 by up to 44% (average 23%) NDCG@10 on Text Retrieval Conference (TREC) Deep Learning Tracks and up to 24% (8% on average) on BEIR.



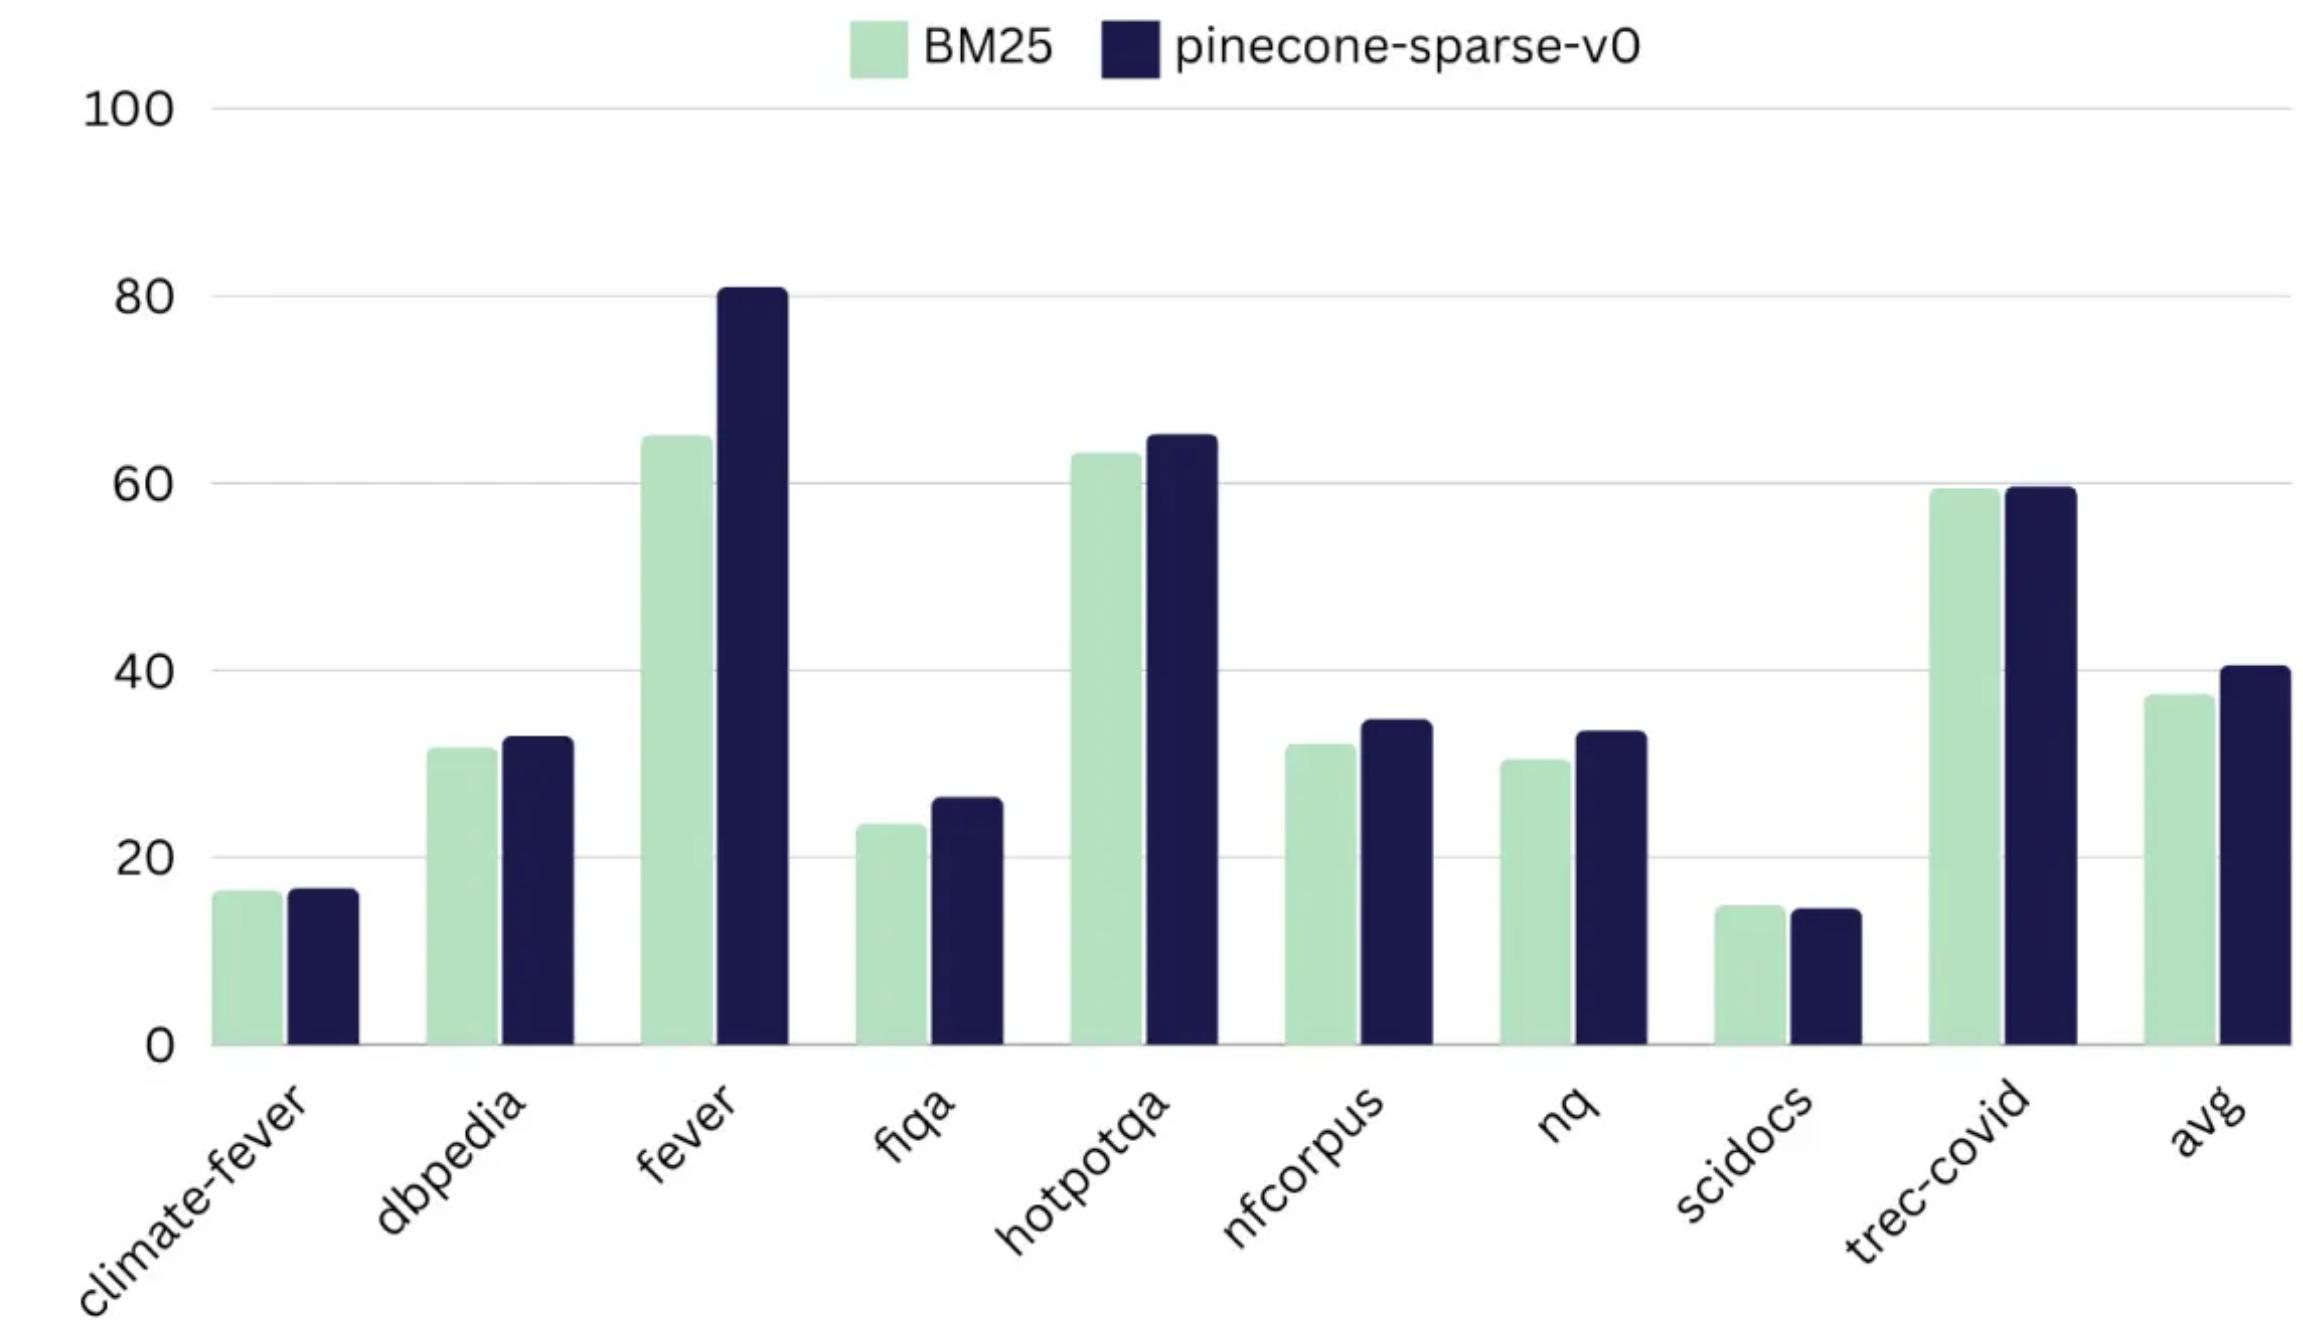

pinecone-sparse-english-v0
PINECONE

Sparse vector model for keyword-style search.

Task	Embedding
Modality	Text
Max Input Tokens	512
Price	\$0.08 / million tokens

[Try this model](#)

Benchmarking pinecone-sparse-english-v0



Contextual Importance

It leverages a combination of a **traditional inverted indexes** and **contextualized language models** for efficient retrieval

It estimates the semantic importance to produces a single value impact score for each tokens of a document collection

pinecone-sparse-v0	BM25	Example
2.46	0.61	Does are the females in the deer family of mammals, individually called a doe.
0.20	0.62	Does anyone else want to come to the movies with us?

Whole-word tokenization

“Top cardiologists in Palo Alto for mitral valve repair”

WordPiece tokenization

```
● ● ●  
{  
  "mit": 1.73,  
  "alto": 1.7,  
  "palo": 1.63,  
  "valve": 1.52,  
  "repair": 1.29,  
  "hospital": 1.15,  
  "##ral": 1.03,  
  "heart": 1.03,  
  "doctor": 1.02,  
  "##ogist": 1.02,  
  "##logist": 0.92,  
  "top": 0.92,  
  "clinic": 0.87,  
  "surgery": 0.76,  
  "card": 0.74,  
  "surgeon": 0.68,  
  "valves": 0.67,  
  "dentist": 0.59,  
  "##iol": 0.59,  
  "clara": 0.53,  
  "specialist": 0.52  
}
```

Not limited to predefined tokens
(e.g. ~32k BERT vocabulary)


No undesired fragmentation of keywords

No addition of unrelated/harmful tokens

Whole-word tokenization


```
● ● ●  
{  
  "mitral": 4.95,  
  "cardiologists": 4.55,  
  "palo": 4.24,  
  "alto": 3.78,  
  "valve": 3.29,  
  "repair": 2.31,  
  "top": 1.33,  
  "in": 0.8,  
  "for": 0.76  
}
```

Query-specific Retrieval Quality




Combining Sources Improves Recall

MS Marco v1 - TREC 2019 & 2020



MS Marco v2 - TREC 2021 & 2022 & 2023



Learned Sparse Retrieval with Entities



Method	Reg	TREC Robust04			TREC Core 2018			CODEC		
		nDCG@10	nDCG@20	R@1k	nDCG@10	nDCG@20	R@1k	nDCG@10	nDCG@20	R@1k
<i>Unsupervised sparse retrieval</i>										
BM25		39.71	36.25	57.18	30.94	29.19	52.19	37.70	35.28	61.25
BM25 + RM3		43.77	40.64	64.21	35.82	34.79	60.09	39.93	39.96	65.70
<i>Zero-shot Dense Retrieval</i>										
DistilBERT-dot-v5		37.95	34.97	52.41	37.02	34.60	54.07	42.76	46.67	60.33
GTR-T5-base		43.79	39.33	54.35	38.81	36.51	57.62	48.42	54.01	66.96
Sentence-T5-base		44.06	39.60	57.64	43.18	39.54	60.88	44.22	32.10	65.48
<i>Learned Sparse Retrieval</i>										
LSR-w	$1e^{-3}$	40.37	37.23	55.66	34.50	31.45	52.66	39.10	35.32	57.58
DyVo (REL)		41.52	38.62	56.78	37.50	34.61	54.14	42.67	38.32	59.81
LSR-w	$1e^{-4}$	47.69	44.48	64.47	38.94	37.37	60.44	50.54	46.71	66.39
DyVo (REL)		48.15	44.85	64.72	43.10	39.46	60.43	51.66	47.95	68.49
LSR-w	$1e^{-5}$	49.13	46.34	66.86	40.99	38.73	63.22	52.61	49.22	69.07
DyVo (REL)		51.19	47.65	68.56	43.72	40.56	63.56	53.40	51.15	70.60

Table 1: Results with linked entities. All LSR models use a DistilBERT backbone. DyVo uses entities found by the REL entity linker and LaQue entity embeddings. All documents are truncated to the first 512 tokens.

Rethinking Sparse Retrieval in the Age of RAG

✓ Treat sparse as a first-class citizen in RAG — not just a fallback for coverage.

🔬 Explore modern sparse models like DeeplImpact, uniCOIL, and SPLADE that combine efficiency with semantic awareness.

💡 Use sparse strategically for:


🔗 Evidence grounding (exact terms matter)

🔄 Multi-hop and graph-based reasoning

⚡ Low-latency, high-throughput deployments

⚖ Evaluate sparse and dense under the same conditions — scale, tuning, budget.

💡 Design hybrid pipelines that do more than fuse — they collaborate.



“RAG may be a generation task – but it’s still retrieval-augmented. Let’s make that retrieval smarter, faster, and sparser.”

Thanks!

Any questions?